

Investigate_a_Dataset_TMDB

May 31, 2022

1 Project: Investigate a Dataset - [TMDB Movie Data]

1.1 Table of Contents

Introduction

Data Wrangling

Exploratory Data Analysis

Conclusions

Introduction

1.1.1 Dataset Description

This dataset contains information about 10,000 movies collected from The Movie Database (TMDB), including user ratings and revenue.

The columns of the dataset are listed below with their respective significance.

id - unique id of the movie

imdb_id - id of the movie by IMDB

popularity - popularity score of the movie

budget - budget used to make the movie

revenue - revenue brought in by the movie

original_title - movie title

cast - list/names of actors in the movie

homepage - website of the movie

director - person who directed the movie

tagline - movie's advertising slogan

keywords - words to easily find a movie title

overview - summary of movie storyline

runtime - total time of the movie

genres - categorization of movie based on plot, story etc.

production_companies - production company of movie

release_date - date of release

vote_count - total number of votes

vote_average - average rating of votes

release_year - year of movie release

budget_adj - budget adjustment due to inflation

revenue_adj - revenue adjustment due to inflation

1.1.2 Questions for Analysis

The following questions will be explored and analysed in this report.

1. Are very popular movies high revenue/grossing movies?
2. Do movies with high budgets get the highest revenues?
3. Do lower user ratings translate to low revenues for movies?
4. How has movie revenue changed over the years?
5. How has movie popularity changed over the years?
6. How has movie budget evolved over the years?
7. Which years were most profitable for movies?

```
In [1]: # Import all necessary packages (Numpy, Pandas, Matplotlib and Seaborn)
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
% matplotlib inline
```

```
In [10]: # Upgrade pandas to use dataframe.explode() function.
!pip install --upgrade pandas==0.25.0
```

Collecting pandas==0.25.0

Downloading <https://files.pythonhosted.org/packages/1d/9a/7eb9952f4b4d73fbd75ad1d5d6112f407e69>

100% || 10.5MB 1.6MB/s eta 0:00:01 12% | | 1.3MB 22.2MB/s eta 0

Requirement already satisfied, skipping upgrade: pytz>=2017.2 in /opt/conda/lib/python3.6/site-p

Collecting numpy>=1.13.3 (from pandas==0.25.0)

Downloading <https://files.pythonhosted.org/packages/45/b2/6c7545bb7a38754d63048c7696804a0d9473>

100% || 13.4MB 2.8MB/s eta 0:00:01 25% | | 3.4MB 18.0MB/s eta 0:00:0

Requirement already satisfied, skipping upgrade: python-dateutil>=2.6.1 in /opt/conda/lib/python

Requirement already satisfied, skipping upgrade: six>=1.5 in /opt/conda/lib/python3.6/site-packa

tensorflow 1.3.0 requires tensorflow-tensorboard<0.2.0,>=0.1.0, which is not installed.

Installing collected packages: numpy, pandas

Found existing installation: numpy 1.12.1

Uninstalling numpy-1.12.1:

Successfully uninstalled numpy-1.12.1

Found existing installation: pandas 0.23.3

Uninstalling pandas-0.23.3:

Successfully uninstalled pandas-0.23.3

Successfully installed numpy-1.19.5 pandas-0.25.0

```
## Data Wrangling
```

1.1.3 General Properties

```
In [2]: #load TMDB movies dataset
df = pd.read_csv('tmdb_movies.csv')
```

```
#inspect few lines
```

```
df.head()
```

```
Out[2]:
```

	id	imdb_id	popularity	budget	revenue	\
0	135397	tt0369610	32.985763	150000000	1513528810	
1	76341	tt1392190	28.419936	150000000	378436354	
2	262500	tt2908446	13.112507	110000000	295238201	
3	140607	tt2488496	11.173104	200000000	2068178225	
4	168259	tt2820852	9.335014	190000000	1506249360	

	original_title	\
0	Jurassic World	
1	Mad Max: Fury Road	
2	Insurgent	
3	Star Wars: The Force Awakens	
4	Furious 7	

	cast	\
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	
2	Shailene Woodley Theo James Kate Winslet Ansel...	
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...	
4	Vin Diesel Paul Walker Jason Statham Michelle ...	

	homepage	director	\
0	http://www.jurassicworld.com/	Colin Trevorrow	
1	http://www.madmaxmovie.com/	George Miller	
2	http://www.thedivergentseries.movie/#insurgent	Robert Schwentke	
3	http://www.starwars.com/films/star-wars-episod...	J.J. Abrams	
4	http://www.furious7.com/	James Wan	

	tagline	...	\
0	The park is open.	...	
1	What a Lovely Day.	...	
2	One Choice Can Destroy You	...	
3	Every generation has a story.	...	
4	Vengeance Hits Home	...	

	overview	runtime	\
0	Twenty-two years after the events of Jurassic ...	124	
1	An apocalyptic story set in the furthest reach...	120	
2	Beatrice Prior must confront her inner demons ...	119	
3	Thirty years after defeating the Galactic Empi...	136	
4	Deckard Shaw seeks revenge against Dominic Tor...	137	

	genres	\
0	Action Adventure Science Fiction Thriller	

```

1 Action|Adventure|Science Fiction|Thriller
2      Adventure|Science Fiction|Thriller
3 Action|Adventure|Science Fiction|Fantasy
4      Action|Crime|Thriller

```

```

                                production_companies release_date vote_count \
0 Universal Studios|Amblin Entertainment|Legenda...      6/9/15      5562
1 Village Roadshow Pictures|Kennedy Miller Produ...      5/13/15      6185
2 Summit Entertainment|Mandeville Films|Red Wago...      3/18/15      2480
3      Lucasfilm|Truenorth Productions|Bad Robot      12/15/15      5292
4 Universal Pictures|Original Film|Media Rights ...      4/1/15      2947

```

```

      vote_average  release_year    budget_adj    revenue_adj
0           6.5         2015  1.379999e+08  1.392446e+09
1           7.1         2015  1.379999e+08  3.481613e+08
2           6.3         2015  1.012000e+08  2.716190e+08
3           7.5         2015  1.839999e+08  1.902723e+09
4           7.3         2015  1.747999e+08  1.385749e+09

```

[5 rows x 21 columns]

```

In [3]: #Show the size of dataframe (number of rows and columns)
df.shape

```

```

Out[3]: (10866, 21)

```

```

In [4]: #Display information about the dataframe (column labels, data types and number of cells)
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
id                10866 non-null int64
imdb_id           10856 non-null object
popularity        10866 non-null float64
budget            10866 non-null int64
revenue           10866 non-null int64
original_title    10866 non-null object
cast              10790 non-null object
homepage          2936 non-null object
director          10822 non-null object
tagline           8042 non-null object
keywords          9373 non-null object
overview          10862 non-null object
runtime           10866 non-null int64
genres            10843 non-null object
production_companies 9836 non-null object
release_date      10866 non-null object
vote_count        10866 non-null int64

```

```

vote_average          10866 non-null float64
release_year          10866 non-null int64
budget_adj            10866 non-null float64
revenue_adj           10866 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB

```

```

In [5]: #Shows brief statistical summary of features(columns) in dataframe
df.describe()

```

```

Out[5]:

```

	id	popularity	budget	revenue	runtime \
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000

	vote_count	vote_average	release_year	budget_adj	revenue_adj
count	10866.000000	10866.000000	10866.000000	1.086600e+04	1.086600e+04
mean	217.389748	5.974922	2001.322658	1.755104e+07	5.136436e+07
std	575.619058	0.935142	12.812941	3.430616e+07	1.446325e+08
min	10.000000	1.500000	1960.000000	0.000000e+00	0.000000e+00
25%	17.000000	5.400000	1995.000000	0.000000e+00	0.000000e+00
50%	38.000000	6.000000	2006.000000	0.000000e+00	0.000000e+00
75%	145.750000	6.600000	2011.000000	2.085325e+07	3.369710e+07
max	9767.000000	9.200000	2015.000000	4.250000e+08	2.827124e+09

1.1.4 Data Cleaning

In the next few cells, I will be cleaning the data by dropping columns that I believe are not relevant to my analysis and also rows of missing data and converting datatypes where necessary.

```

In [6]: #Drop columns that will not be used in analysis

```

```

df.drop(['id', 'imdb_id', 'cast', 'homepage', 'keywords', 'tagline', 'overview', 'origin

```

```

In [7]: #check shape of dataframe after dropping columns
df.shape

```

```

Out[7]: (10866, 10)

```

```

In [8]: #check dataframe information again
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 10 columns):
popularity      10866 non-null float64
budget          10866 non-null int64
revenue         10866 non-null int64
genres          10843 non-null object
release_date    10866 non-null object
vote_count      10866 non-null int64
vote_average    10866 non-null float64
release_year    10866 non-null int64
budget_adj      10866 non-null float64
revenue_adj     10866 non-null float64
dtypes: float64(4), int64(4), object(2)
memory usage: 849.0+ KB

```

```
In [9]: df.isnull().sum()
```

```

Out[9]: popularity      0
        budget          0
        revenue         0
        genres          23
        release_date     0
        vote_count       0
        vote_average     0
        release_year     0
        budget_adj       0
        revenue_adj      0
        dtype: int64

```

The `df.info()` and `df.isnull().sum()` shows that genres column is missing some data in some rows. Therefore, we will drop those rows too.

```

In [10]: #Drop null rows from dataframe
         df.dropna(inplace=True)

```

```
In [11]: df.isnull().sum().sum()
```

```
Out[11]: 0
```

There are no null values in the dataframe. Let's finally check for duplicates and drop any duplicates we may find.

```
In [12]: sum(df.duplicated())
```

```
Out[12]: 1
```

```
In [13]: df.drop_duplicates(inplace=True)
```

```
In [14]: sum(df.duplicated())
```

```
Out[14]: 0
```

Let's finally run `df.info()` and `df.describe()` again

```
In [15]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10842 entries, 0 to 10865
Data columns (total 10 columns):
popularity      10842 non-null float64
budget          10842 non-null int64
revenue         10842 non-null int64
genres          10842 non-null object
release_date    10842 non-null object
vote_count      10842 non-null int64
vote_average    10842 non-null float64
release_year    10842 non-null int64
budget_adj      10842 non-null float64
revenue_adj     10842 non-null float64
dtypes: float64(4), int64(4), object(2)
memory usage: 931.7+ KB
```

```
In [16]: df.describe()
```

```
Out[16]:
```

	popularity	budget	revenue	vote_count	vote_average \
count	10842.000000	1.084200e+04	1.084200e+04	10842.000000	10842.000000
mean	0.647461	1.465531e+07	3.991138e+07	217.823649	5.974064
std	1.001032	3.093971e+07	1.171179e+08	576.180993	0.934257
min	0.000065	0.000000e+00	0.000000e+00	10.000000	1.500000
25%	0.208210	0.000000e+00	0.000000e+00	17.000000	5.400000
50%	0.384532	0.000000e+00	0.000000e+00	38.000000	6.000000
75%	0.715393	1.500000e+07	2.414118e+07	146.000000	6.600000
max	32.985763	4.250000e+08	2.781506e+09	9767.000000	9.200000

	release_year	budget_adj	revenue_adj
count	10842.000000	1.084200e+04	1.084200e+04
mean	2001.314794	1.758712e+07	5.147797e+07
std	12.813617	3.433437e+07	1.447723e+08
min	1960.000000	0.000000e+00	0.000000e+00
25%	1995.000000	0.000000e+00	0.000000e+00
50%	2006.000000	0.000000e+00	0.000000e+00
75%	2011.000000	2.092507e+07	3.387838e+07
max	2015.000000	4.250000e+08	2.827124e+09

It looks like the min, 25% and 50% are returning zeros which look like errors. it is likely that there are so many zeros for budget and revenue in the dataframe and this is very unlikely for movies. To continue I'll drop all the rows with zero.

```
In [17]: df = df.replace(0, np.nan)
df = df.dropna()
```

Let's run `df.info()` and `df.describe()` one last time to complete our data cleaning process

```
In [18]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3854 entries, 0 to 10848
Data columns (total 10 columns):
popularity      3854 non-null float64
budget          3854 non-null float64
revenue         3854 non-null float64
genres          3854 non-null object
release_date    3854 non-null object
vote_count      3854 non-null int64
vote_average    3854 non-null float64
release_year    3854 non-null int64
budget_adj      3854 non-null float64
revenue_adj     3854 non-null float64
dtypes: float64(6), int64(2), object(2)
memory usage: 331.2+ KB
```

```
In [19]: df.describe()
```

```
Out[19]:
```

	popularity	budget	revenue	vote_count	vote_average \
count	3854.000000	3.854000e+03	3.854000e+03	3854.000000	3854.000000
mean	1.191554	3.720370e+07	1.076866e+08	527.720291	6.168163
std	1.475162	4.220822e+07	1.765393e+08	879.956821	0.794920
min	0.001117	1.000000e+00	2.000000e+00	10.000000	2.200000
25%	0.462368	1.000000e+07	1.360003e+07	71.000000	5.700000
50%	0.797511	2.400000e+07	4.480000e+07	204.000000	6.200000
75%	1.368324	5.000000e+07	1.242125e+08	580.000000	6.700000
max	32.985763	4.250000e+08	2.781506e+09	9767.000000	8.400000

	release_year	budget_adj	revenue_adj
count	3854.000000	3.854000e+03	3.854000e+03
mean	2001.261028	4.423999e+07	1.370647e+08
std	11.282575	4.480925e+07	2.161114e+08
min	1960.000000	9.693980e-01	2.370705e+00
25%	1995.000000	1.309053e+07	1.835735e+07
50%	2004.000000	3.001611e+07	6.173068e+07
75%	2010.000000	6.061307e+07	1.632577e+08
max	2015.000000	4.250000e+08	2.827124e+09

```
In [20]: df.shape
```

```
Out[20]: (3854, 10)
```



```
In [21]: #Resetting the index to make the numbering correct
df = df.reset_index(drop=True)
```

Before we confirm if the index has been reset correctly, **let's create a simple function that will display few values at both ends of the TMDB dataframe**. This can now help us not to run `df.tail()` and `df.head()` individually all the time.

```
In [22]: #This function when called anytime will display both df.head() and df.tail() together.

def df_ends(df, x=5):
    return df.head(x).append(df.tail(x))
```

```
In [24]: #Now let's call our new function 'df_ends' to confirm that the index has been reset cor

df_ends(df,2)
```

```
Out[24]:
```

	popularity	budget	revenue	\
0	32.985763	150000000.0	1.513529e+09	
1	28.419936	150000000.0	3.784364e+08	
3852	0.299911	12000000.0	2.000000e+07	
3853	0.207257	5115000.0	1.200000e+07	

	genres	release_date	vote_count	\
0	Action Adventure Science Fiction Thriller	6/9/15	5562	
1	Action Adventure Science Fiction Thriller	5/13/15	6185	
3852	Action Adventure Drama War Romance	12/20/66	28	
3853	Adventure Science Fiction	8/24/66	42	

	vote_average	release_year	budget_adj	revenue_adj
0	6.5	2015	1.379999e+08	1.392446e+09
1	7.1	2015	1.379999e+08	3.481613e+08
3852	7.0	1966	8.061618e+07	1.343603e+08
3853	6.7	1966	3.436265e+07	8.061618e+07

Yes, our function works great!

Our data cleaning process is complete. All NaN values, duplicates and columns that are not relevant have been dropped. We also dropped rows for budget and revenue which had zeros as it is unlikely that this is possible in real life scenario and we would want to analyse data where those values for said columns are not zero.

Exploratory Data Analysis


1.1.5 Research Question 1 (Which features can be most associated with high revenue?)

Sub questions: 1. Are very popular movies high revenue/grossing movies? 2. Do movies with high budgets get the highest revenues? 3. Do lower user ratings translate to low revenues for movies?

Please note that for the purposes of the exploratory data analysis, the columns 'revenue_adj' and 'budget_adj' will be used as revenue and budget values respectively and not 'revenue' and 'budget' columns.


```
In [25]: # Explore the revenue_adj feature using a histogram and dataframe.describe()

df.revenue_adj.hist(figsize = (8, 5));
```




The revenue_adj histogram appears very skewed to the right
Let's quickly explore the relationship between revenue_adj and a few features using the scatter diagram.

```
In [26]: df.plot(x = 'revenue_adj', y = 'popularity', kind = 'scatter');
```




It looks like popularity for most movies is low regardless of revenue. This will however be explored in more details.

```
In [27]: df.plot(x = 'revenue_adj', y = 'budget_adj', kind = 'scatter');
```



From initial observations, it looks like budget largely has a positive correlation with revenue.

```
In [28]: df.plot(x = 'revenue_adj', y = 'vote_average', kind = 'scatter');
```



From initial observations, it looks like vote_average (user ratings) has a weak correlation with revenue. We will need to do more detailed exploration.

```
In [29]: #Statistics of revenue feature
```

```
df.revenue_adj.describe()
```

```
Out[29]: count    3.854000e+03  
mean      1.370647e+08  
std       2.161114e+08  
min       2.370705e+00  
25%      1.835735e+07  
50%      6.173068e+07  
75%      1.632577e+08  
max       2.827124e+09  
Name: revenue_adj, dtype: float64
```

Level of movie revenue:

High: 75% to max

Moderately High: 50% to 75%

Medium: 25% to 50%

Low: min to 25%

```
In [30]: # Cutting the data from revenue_adj into groups from the describe results
        bin_edges = [2.370705e+00, 1.835735e+07, 6.173068e+07, 1.632577e+08, 2.827124e+09]
```

We are grouping the revenue_adj data to create a new feature.

```
In [31]: # We create labels for the four different revenue levels and name each revenue level
        bin_names = ['low', 'medium', 'moderately high', 'high']
```

The new feature will need new values as found in the cell above.

```
In [32]: # Creates revenue_levels column
        df['revenue_levels'] = pd.cut(df['revenue_adj'], bin_edges, labels=bin_names)
```

We have successfully created a new column called 'revenue_levels' where 'revenue_adj' values have been categorised in terms of high, moderately high, medium and low.

```
In [35]: #Let's check if the revenue_levels column is showing
```

```
df_ends(df,2)
```

```
Out[35]:
```

	popularity	budget	revenue \
0	32.985763	150000000.0	1.513529e+09
1	28.419936	150000000.0	3.784364e+08
3852	0.299911	12000000.0	2.000000e+07
3853	0.207257	5115000.0	1.200000e+07

	genres	release_date	vote_count \
0	Action Adventure Science Fiction Thriller	6/9/15	5562
1	Action Adventure Science Fiction Thriller	5/13/15	6185
3852	Action Adventure Drama War Romance	12/20/66	28
3853	Adventure Science Fiction	8/24/66	42

	vote_average	release_year	budget_adj	revenue_adj	revenue_levels
0	6.5	2015	1.379999e+08	1.392446e+09	high
1	7.1	2015	1.379999e+08	3.481613e+08	high
3852	7.0	1966	8.061618e+07	1.343603e+08	moderately high
3853	6.7	1966	3.436265e+07	8.061618e+07	moderately high

Now let's use groupby to do more exploration of revenue levels with the following features; vote_average, popularity and budget_adj

```
In [36]: # So we will first find the mean popularity of each revenue level with groupby
        df.groupby('revenue_levels')['popularity'].mean()
```

```
Out[36]: revenue_levels
low                0.559361
medium             0.808519
moderately high    1.148874
high               2.249021
Name: popularity, dtype: float64
```

We see that the level of revenue that receives the highest mean popularity score is 'high'

```
In [37]: # Next, we will find the mean user rating/vote_average of each revenue level with groupby
df.groupby('revenue_levels')['vote_average'].mean()
```

```
Out[37]: revenue_levels
low                5.958299
medium             6.066355
moderately high    6.188681
high               6.459232
Name: vote_average, dtype: float64
```

We see that the level of revenue that receives the highest mean vote average is 'high'

```
In [38]: # Finally, we will find the mean budget of each revenue level with groupby
df.groupby('revenue_levels')['budget_adj'].mean()
```

```
Out[38]: revenue_levels
low                1.599161e+07
medium             2.947073e+07
moderately high    4.449050e+07
high               8.699208e+07
Name: budget_adj, dtype: float64
```

Again, we see that the level of revenue with the highest mean budget is 'high'

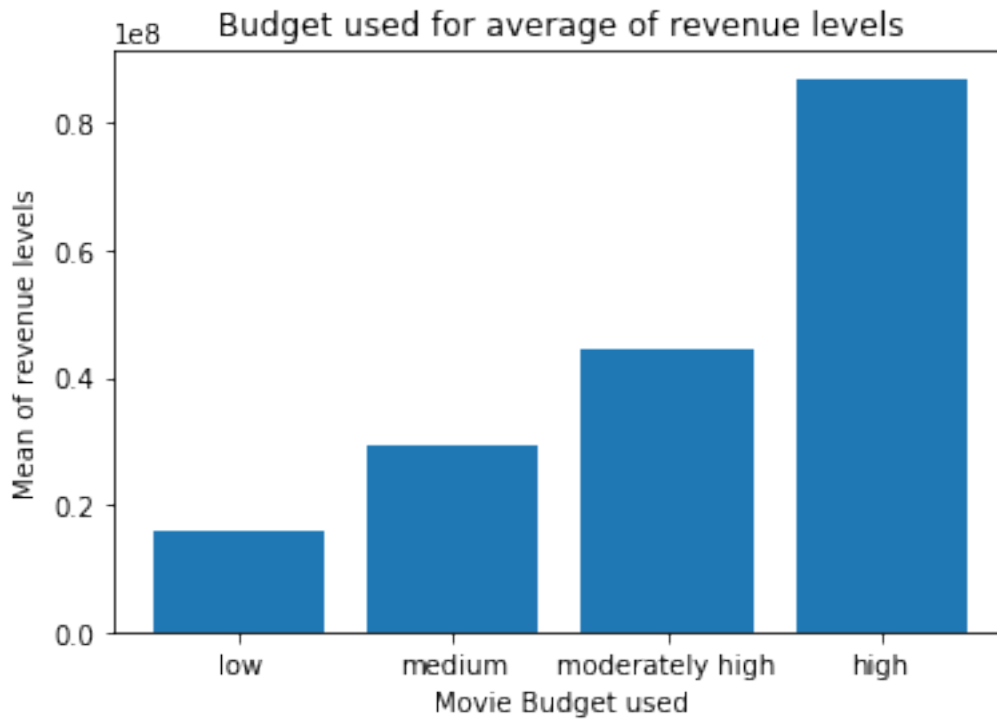
Let's explore both revenue level means and budget means, popularity means and vote average means again using Matplotlib visualisation.

Let's create more plots using Matplotlib

```
In [39]: revenue_levels_means = df.groupby('revenue_levels')['budget_adj'].mean()
```

```
In [40]: # Create a bar chart with proper labels
locations = [1, 2, 3, 4]
heights = revenue_levels_means
labels = ['low', 'medium', 'moderately high', 'high']
plt.bar(locations, heights, tick_label=labels)
plt.title('Budget used for average of revenue levels')
plt.xlabel('Movie Budget used')
plt.ylabel('Mean of revenue levels')
```

```
Out[40]: Text(0,0.5,'Mean of revenue levels')
```

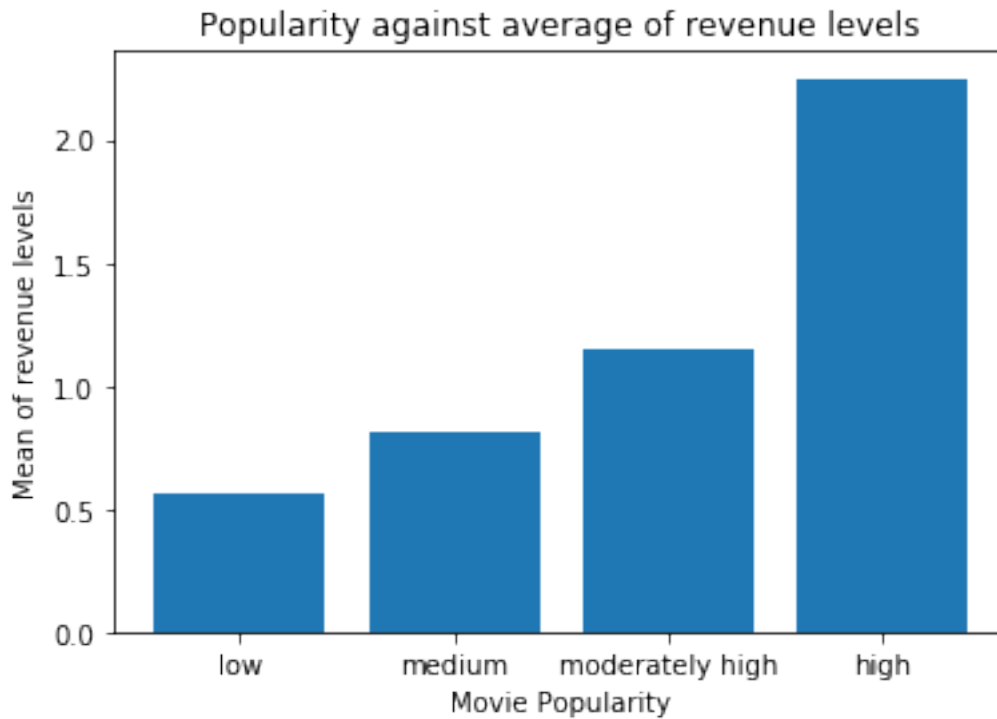


From the plot above, we see that averagely higher budget movies are associated with higher revenue levels.

```
In [41]: revenue_levels_means1 = df.groupby('revenue_levels')['popularity'].mean()
```

```
In [42]: # Create a bar chart with proper labels
locations = [1, 2, 3, 4]
heights = revenue_levels_means1
labels = ['low', 'medium', 'moderately high', 'high']
plt.bar(locations, heights, tick_label=labels)
plt.title('Popularity against average of revenue levels')
plt.xlabel('Movie Popularity')
plt.ylabel('Mean of revenue levels')
```

```
Out[42]: Text(0,0.5,'Mean of revenue levels')
```

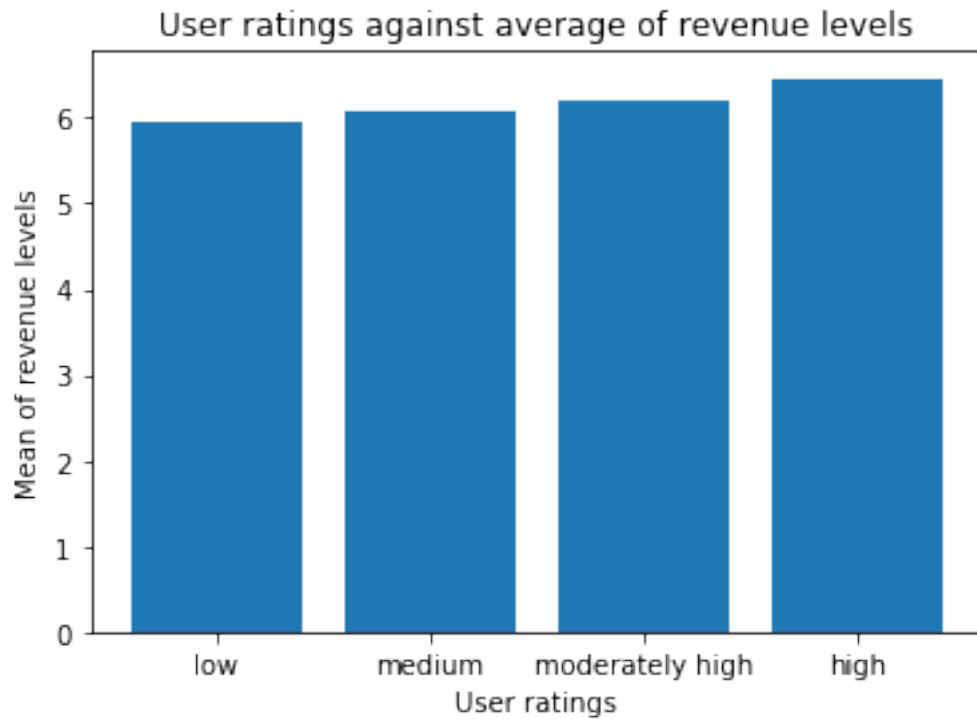


We observe from the plot above that averagely high revenue level movies seem to be associated with high popularity.

```
In [43]: revenue_levels_means2 = df.groupby('revenue_levels')['vote_average'].mean()
```

```
In [44]: # Create a bar chart with proper labels
locations = [1, 2, 3, 4]
heights = revenue_levels_means2
labels = ['low', 'medium', 'moderately high', 'high']
plt.bar(locations, heights, tick_label=labels)
plt.title('User ratings against average of revenue levels')
plt.xlabel('User ratings')
plt.ylabel('Mean of revenue levels')
```

```
Out[44]: Text(0,0.5,'Mean of revenue levels')
```

We observe that the revenue levels doesn't seem to be affected by user ratings.
This marks the end of the first major part of our data exploration.


1.1.6 Research Question 2 (How has movie success changed over the years?)

Sub questions:

1. How has movie revenue changed over the years?
2. How has movie budgets changed over the years?
3. How has popularity changed over the years?
4. Which years were most profitable for movies?


We shall use histograms, scatter plots and bar charts to investigate the way movies have changed over the years in terms of budget, revenue, profit and popularity.

```
In [45]: df.release_year.hist(figsize = (8, 5));
```




Firstly, we observe that the release year histogram is skewed to the left.
Let's create some scatter plots related to our questions

```
In [46]: df.plot(x = 'release_year', y = 'revenue_adj', kind = 'scatter');
```




From the scatter plot it looks like revenue has increased overall between 1960s and 2010s. However, there were some individual years where revenue was very high.

```
In [47]: df.plot(x = 'release_year', y = 'budget_adj', kind = 'scatter');
```



From the scatter plot it looks like movie budget has also increased overall between 1960s and 2010s.

```
In [48]: df.plot(x = 'release_year', y = 'popularity', kind = 'scatter');
```



The scatter plot shows there is a positive correlation between popularity and the increase of years, although it appears the trend appears slow.

We have about 56 years of movie data. It will be much better to investigate the long term change over the years in terms of decades. For this we will need to group the years into decades and create a new column in the dataframe called 'decades'.

```
In [49]: # Group/'cut' the years into decades
         bin_edges = [1960, 1970, 1980, 1990, 2000, 2010, 2015]

In [50]: # Label the new decades
         bin_names = ['1960s', '1970s', '1980s', '1990s', '2000s', '2010s']

In [51]: # Create the new 'decade' column
         df['decade'] = pd.cut(df['release_year'], bin_edges, labels=bin_names)

In [52]: df.head(2)
```

	popularity	budget	revenue \
0	32.985763	150000000.0	1.513529e+09
1	28.419936	150000000.0	3.784364e+08

	genres	release_date	vote_count \
0	Action Adventure Science Fiction Thriller	6/9/15	5562
1	Action Adventure Science Fiction Thriller	5/13/15	6185

	vote_average	release_year	budget_adj	revenue_adj	revenue_levels \
--	--------------	--------------	------------	-------------	------------------

0	6.5	2015	1.379999e+08	1.392446e+09	high
1	7.1	2015	1.379999e+08	3.481613e+08	high

decade	
0	2010s
1	2010s

Last two cells are for checking if our new column was successfully created.
 Let's plot a bar chart using Matplotlib to show correlation between Decades and Revenue


```
In [54]: df.groupby('decade')['revenue_adj'].mean()
```

```
Out[54]: decade
1960s    2.842694e+08
1970s    2.567380e+08
1980s    1.274362e+08
1990s    1.380226e+08
2000s    1.225630e+08
2010s    1.330990e+08
Name: revenue_adj, dtype: float64
```

```
In [55]: # We plot the bar chart to explore the mean revenue generated in the various decades
revenue_during_decade = df.groupby('decade')['revenue_adj'].mean()

plt.bar(revenue_during_decade.index, revenue_during_decade.values)
plt.title('Movie Revenue over the Years')
plt.xlabel('Decade')
plt.ylabel('Revenue')
```

```
Out[55]: Text(0,0.5,'Revenue')
```



From the bar chart above, we see that average movie revenue was highest in the 1960s and 1970s and has steadily come down to half of that over the rest of the decades.

Let's plot a bar chart using Matplotlib to show correlation between Decades and Mean Budget used


```
In [56]: df.groupby('decade')['budget_adj'].mean()
```

```
Out[56]: decade
1960s    4.546152e+07
1970s    2.851827e+07
1980s    2.957904e+07
1990s    4.988054e+07
2000s    4.766348e+07
2010s    4.271093e+07
Name: budget_adj, dtype: float64
```

```
In [57]: # We plot the bar chart to explore the mean budget used over the various decades
budget_during_decade = df.groupby('decade')['budget_adj'].mean()

plt.bar(budget_during_decade.index, budget_during_decade.values)
plt.title('Mean Movie Budget over the Years')
plt.xlabel('Decade')
plt.ylabel('Average Budget')
```

```
Out[57]: Text(0,0.5,'Average Budget')
```



From the bar chart, we see that budget used has high in the 1960s, then it came low through the 1970s and 1980s and eventually went up in 1990s and is still high till 2010s although there is a very small decline between 1990s and 2010s.

Let's plot a bar chart using Matplotlib to show Average Popularity of movies over the Decades


```
In [58]: df.groupby('decade')['popularity'].mean()
```

```
Out[58]: decade
1960s    0.900749
1970s    1.044932
1980s    0.848137
1990s    0.938220
2000s    1.082989
2010s    1.841155
Name: popularity, dtype: float64
```

```
In [59]: # We plot the bar chart to explore the mean budget used over the various decades
popularity_during_decade = df.groupby('decade')['popularity'].mean()

plt.bar(popularity_during_decade.index, popularity_during_decade.values)
plt.title('Average Movie Popularity over the Years')
plt.xlabel('Decade')
plt.ylabel('Average Movie Popularity')
```

```
Out[59]: Text(0,0.5,'Average Movie Popularity')
```



We see from the chart above that movies have relatively maintained their popularity over the decades until the 2010s where they seem to have gone up.

One more way we want to explore our data over the decades is in terms of profits made from the movies. We define profit here as 'revenue_adj - budget_adj'.

Let's create a new column in the dataframe called 'profit'

```
In [60]: df['profit'] = df['revenue_adj'] - df['budget_adj']
```

```
In [62]: #Let's confirm that the 'profit' column has been created.
df_ends(df,2)
```

```
Out[62]:
```

	popularity	budget	revenue \
0	32.985763	150000000.0	1.513529e+09
1	28.419936	150000000.0	3.784364e+08
3852	0.299911	12000000.0	2.000000e+07
3853	0.207257	5115000.0	1.200000e+07

	genres	release_date	vote_count \
0	Action Adventure Science Fiction Thriller	6/9/15	5562
1	Action Adventure Science Fiction Thriller	5/13/15	6185
3852	Action Adventure Drama War Romance	12/20/66	28
3853	Adventure Science Fiction	8/24/66	42

	vote_average	release_year	budget_adj	revenue_adj	revenue_levels \
0	6.5	2015	1.379999e+08	1.392446e+09	high

1	7.1	2015	1.379999e+08	3.481613e+08	high
3852	7.0	1966	8.061618e+07	1.343603e+08	moderately high
3853	6.7	1966	3.436265e+07	8.061618e+07	moderately high

	decade	profit
0	2010s	1.254446e+09
1	2010s	2.101614e+08
3852	1960s	5.374412e+07
3853	1960s	4.625353e+07

Let's plot a bar chart using Matplotlib to show Average Popularity of movies over the Decades


```
In [63]: df.groupby('decade')['profit'].mean()
```

```
Out[63]: decade
1960s    2.388079e+08
1970s    2.282197e+08
1980s    9.785714e+07
1990s    8.814204e+07
2000s    7.489950e+07
2010s    9.038803e+07
Name: profit, dtype: float64
```

```
In [64]: # We plot the bar chart to see the average profit made over the various decades
profit_during_decade = df.groupby('decade')['profit'].mean()

plt.bar(profit_during_decade.index, profit_during_decade.values)
plt.title('Average Movie Profit made over the Years')
plt.xlabel('Decade')
plt.ylabel('Average Profit')
```

```
Out[64]: Text(0,0.5,'Average Profit')
```



From the plot above, the 1960s and 1970s look to have been the most profitable years for movies. The subsequent decades show even less than half of that average profit.

Conclusions

Main question 1: Which features can be most associated with high revenue? (1 to 3)

1. Are very popular movies high revenue/grossing movies?

It can be deduced from the exploratory analysis that on average, higher revenue levels can be associated with slightly higher popularity levels.

2. Do movies with high budgets get the highest revenues?

From the data exploration, we see that averagely, higher budget movies are associated with higher revenue levels.

3. Do lower user ratings translate to low revenues for movies?

It can be inferred from the exploration of data that user ratings (vote average) does not directly influence the revenue levels of movies.

Main question 2: How has movie success changed over the years? (4 to 7)

4. How has movie revenue changed over the years?

The bar chart of average movie revenue for the various decades shows that average movie revenue was highest in the 1960s and 1970s and steadily came down to half of that over the rest of the decades.

5. How has movie popularity changed over the years?

The exploratory analysis shows that movies have relatively maintained their popularity level over the decades until the 2010s where the popularity appears to have doubled.

6. How has movie budget evolved over the years?

From the analysis, movie budget was high in the 1960s, then it came low through the 1970s and 1980s and eventually went up in 1990s and is still high till 2010s although there is a very small decline between 1990s and 2010s.

7. Which years were most profitable for movies?

The analysis shows that the 1960s and 1970s look to have been the most profitable years for movies. The subsequent decades show even less than half of that average profit.

Additional comments under conclusion > The genre column can be explored to investigate its relationship with revenue levels as well as its popularity over the decades.

1.1.7 Limitations

1. The data had a lot of zero values and null values which resulted in deletion of many rows. Out of the initial 10866 rows of data, only 3854 was used in the analysis after cleaning. This may have resulted in loss of valuable data.
2. Although, budget_adj and revenue_adj were two of the main columns used in the dataframe for analysis, the currencies of the values were not added. This means that the values may not be very accurate for analysis.

1.2 Reference

1. https://pandas.pydata.org/docs/user_guide/index.html#user-guide
2. <https://stackoverflow.com/questions/22649693/drop-rows-with-all-zeros-in-pandas-data-frame>
3. <https://classroom.udacity.com>

Extra note: This updated .ipynb file has the specified change recommended from the 1st review, which is the creation of a function.

```
In [66]: from subprocess import call
         call(['python', '-m', 'nbconvert', 'Investigate_a_Dataset_TMDB.ipynb'])
```

```
Out[66]: 0
```

```
In [ ]:
```